R tools for accessing research literature for text mining

Scott Chamberlain*,[a]

[a]*rOpenSci, Museum of Paleontology, University of California, Berkeley, CA, USA*

Abstract

Text mining is a powerful method for answering research questions. However, getting texts to extract information can be a daunting and complicated task. The primary reason for this is the diversity of publisher technologies. There are thousands of different publishers, each with their own licenses, URL patterns, access options, and more. Layered on top of that is the varied access each user has based on their institutional affiliation. Here, I introduce a suite of software packages in the R programming language for fetching texts. The tapestry of different publishers, access levels, and other factors requires a patchwork of approaches for getting texts to users. The flagship R package called fulltext attempts to simplify search and retrieval of texts for text mining by serving as an interface to the varied and complex publishers. The fulltext package, along with many others, make acquiring texts easier than ever, facilitating answering research questions with text mining.

*Corresponding author

Email address: myrmecocystus(at)gmail.com (Scott Chamberlain)

August 22, 2019

Introduction

There's more than 100 million research articles published (Crossref API: https://github.com/CrossRef/rest-api-doc), representing an enormous amount of knowledge. In addition to simply reading these articles, the articles contain a vast trove of information of interest to researchers for machine aided questions (Kong & Gerstein, 2018; Usai et al., 2018). For example, many researchers are interested in statistical outcomes of articles that can be extracted from numeric results: P-values, effect sizes, means, and more. In addition, researchers are often interested in words in articles, their use through time, and the contexts they are found in.

Text mining is the broad term associated with pulling information out of articles. Given the importance of text mining, good text mining tools are needed to make it easier for researchers to do. Graphical user interface (GUI) based text mining tools are available (e.g., Ba & Bossy, 2016; Cañada et al., 2017; Muñoz, Kissling & Loon, 2019) and some research papers have used them (Chaix et al., 2019), but given the urgent recent call to action for more reproducible research (Open Science Collaboration, 2015; Camerer et al., 2016, 2018), we must move away from GUI based tools as fast as possible. A number of examples of programmatic tools can be found in the literature. For example, Sinclair et al. (2016) present a tool in Python called seqenv for the domain specific task of linking sequences to environments through text mining.

Most recent text mining papers do not use programmatic approaches, highlighting the need for more programmatic text mining tools, and increased discussion of those tools to increase awareness. For example, many papers search Web of Science using their web interface, and downloading papers manually (Ding, Li & Fan, 2018; McCallen et al., 2019). Many of these papers doing GUI based searching and paper downloads are using R or Python downstream for analysis; replacing GUI based data acquisition with programmatic approaches will improve research.

The R programming language is free of cost, and is used widely throughout many academic fields; tools in R for text mining are of particular importance because they can be adopted by academics rapidly.

Here, I present an overview of text mining tools in the R programming language, not for text mining analysis, but rather those tools for searching for, acquiring, and extracting parts of texts (e.g., title, abstract, authors). Most of the packages presented here are part of the rOpenSci suite (https://ropensci.org/).

Digital articles: technical aspects

Those articles that are digital (which in theory includes all articles) can be split into two groups: machine readable and non-machine readable.

The machine readable articles are those in XML[1], JSON[2], or plain text format. The former two, XML and JSON, are the best machine readable types because they are structured data[3], whereas plain text has no structure - it's simply a set of characters with line breaks and spaces in between.

Of the non-machine readable types, the most noteable is the Portable Document Format (PDF)[4]. These can be broken out into two groups: text based PDFs and scanned PDFs. The former are converted from digital versions of various kinds (MS Word, OpenOffice, LaTeX, markdown, etc.), while the latter are created by scanning print articles to a PDF format. Text-based PDFs are much better for text mining purposes as plain text can be extracted easily in R with pdftools, a binding to libpoppler. However, with scanned PDFs, text must be extracted using Optimal Character Recognition (OCR; see R package tesseract), which isn't always a clean solution, especially compared to true text based PDFs.

The reality in scholarly publishing is all publishers, if they provide any access to their articles, only provide PDF format. Very few publishers, with some quite large (Elsevier, Pensoft, PLOS), provide XML format. Although most publishers most likely have the XML behind each of their articles, they for some indefensible reason do not share it - making text mining more difficult. Some provide plain text (Elsevier). I only know of one publisher that provides full text as JSON (PLOS). Thus, text mining, in most cases, will require extracting text from PDFs.

Digital articles: the access landscape

Acces to full-text is the holy grail in text mining. Some use cases can get by with article metadata (authors, title, etc.), some with abstracts, but many use cases require full-text.

The landscape of access to full-text is extremely hetergeous, with the majority of variation along the publisher axis. The major hurdle is paywalls. The majority of articles are published by the big three publishers - Wiley, Springer, Elsevier - and the majority of their articles are behind paywalls.

[1] https://www.w3.org/TR/xml/
[2] https://tools.ietf.org/html/rfc7159
[3] https://en.wikipedia.org/wiki/Data_model
[4] https://en.wikipedia.org/wiki/PDF

A promising sign is an increasing number of open access articles, yet open access articles represent a small percent of all articles: an estimate in 2018 said that 28% of the scholarly literature was open access (Piwowar et al., 2018).

With respect to paywalled articles, access varies by institution, depending on each institution's publisher contracts. MORE ABOUT THIS ...

Some may not realize access to articles varies with IP address so that access from campus vs. from home (if not on a VPN) will drastically differ. Sometimes a VPN is required, and this can provide a significant technical hurdle to users attempting to do text mining work.

One final hurdle in text mining comes unsurprisingly from Elsevier. They use so-called "fences" for programmatic access. That is, even if a person trying to get an article programmatically their institution has access to and they have access to, and they are on the correct IP address, they may still not get access to an Elsevier article. Elsevier puts in place these fences and only if you contact their technical team directly can you get these fences removed, and only then on a per institution basis.

I can not end this section without mentioning SciHub. This is a last resort option for many probably (or possibly first, depending on your level of access), providing access to full text of articles that are normally paywalled. No tools in this manuscript provide access to SciHub.

The discovery problem

A text mining project starts with a question. From that question, researchers then attempt to acquire scholarly articles for text mining. Finding appropriate articles is not altogether straight-forward.

Some of the discovery difficulty relates to the fact that there are so many places to search for articles; a non-exhaustive list: Google Scholar, Microsoft Academic Research, Scopus, ScienceDirect, Web of Science, Pubmed/Entrez, Europe PMC, Directory of Open Access Journals, Open Knowledge Maps, CORE, Fatcat, and more. It's probably difficult to know where the best place is to search. Some of these are paywalled (e.g., Web of Science), and some are not.

The most important aspect about any source for article search with respect to reproducible research is being able to use the data source programmatically. Of those listed above, the following can be used programmatically: Microsoft Academic Research, Scopus, ScienceDirect, Pubmed/Entrez, Europe PMC, and Directory of Open Access Journals. All of these are included in the R package fulltext, discussed further below.

On top of the vast array of different data sources is the varied ways that search is implemented in each source. Most sources are probably using Solr or Elasticsearch under the hood, though we can't know this for sure as most do not make their software infrastructure public knowledge. Nonetheless, data sources differ in how search works from the user perspective. For example, some provide wild card search and some do not. Some sources are searching full text of articles, while others only search metadata (i.e., title, authors, abstract). In addition, each source has a different set of metadata/full text available. In brief, the same search against different sources produces different results. Some text mining research articles perform the same search against many different sources (refs), while others choose just one source.

Data sources

There is increasing open access scientific literature content available online. However, only a small proportion of scientific journals provide access to their full content; whereas, most publishers provide open access to their metadata only (most often through Crossref; Table 1). The following is a synopsis of the major data sources and associated R tools.

Table 1. Sources of scientific literature, their content type provided via web services, whether rOpenSci has an R packages for the service, and where to find the API documentation.

Data Provider	Content Type	rOpenSci Package	Documentation
Crossref	Metadata	rcrossref/crminer	[5]
DataCite	Metadata	rdatacite	[6]
Biodiversity Heritage Library	Full content/Metadata	rbhl	[7]
Public Library of Science (PLoS)	Full content/altmetrics	rplos	[8]
Scopus (Elsevier)	Full content/Metadata	fulltext	[9]
arXiv	Full content/Metadata	aRxiv	[10]
Biomed Central (via Springer)	Full content/Metadata	fulltext	[11]
bioRxiv	Full content/Metadata	fulltext	[12]
PMC/Pubmed (via Entrez)	Full content/Metadata	rentrez	[13]
Europe PMC	Full content/Metadata	europepmc	[14]
Microsoft Academic Search	Metadata	microdemic	[15]
Directory of Open Access Journals	Metadata	jaod	[16]
JSTOR Data for Research	Full content	jstor	[17]
ORCID	Metadata	rorcid	[18]
Wikimedia's Citoid	Citations	rcitoid	[19]
Open Citation Corpus	Citations	citecorp	[20]

[5] https://api.crossref.org
[6] https://support.datacite.org/docs/api
[7] http://bit.ly/KYQ1Rd
[8] http://api.plos.org/solr
[9] http://bit.ly/J9S616
[10] https://arxiv.org/help/api/index
[11] https://dev.springer.com/
[12] http://www.biorxiv.org/
[13] https://www.ncbi.nlm.nih.gov/books/NBK25500
[14] https://azure.microsoft.com/en-us/services/cognitive-services
[15] https://dev.labs.cognitive.microsoft.com/docs/services/56332331778daf02acc0a50b/operations/565d9001ca73072048922d97
[16] https://doaj.org/api/v1/docs
[17] https://www.jstor.org/dfr/
[18] https://pub.orcid.org/
[19] https://en.wikipedia.org/api/rest_v1/#/Citation/getCitation
[20] http://opencitations.net/

Data Provider	Content Type	rOpenSci Package	Documentation
Fatcat	Metadata	none	[21]
SHERPA/RoMEO	Journal Level Metadata	rromeo	[22]
CORE	Full content/Metadata	rcoreoa	[23]
Dissemin	Metadata	dissemr	[24]

Crossref/Datacite

Crossref is a non-profit that creates (or "mints") Digital Object Identifiers (DOIs). In addition, they maintain metadata associated with each DOI. The metadata ranges from simple (including author, title, dates, DOI, type, publisher) to including number of citations to the article, as well as references in the article, and even abstracts. At the time of writing they hold 100 million DOIs.

One can search by DOI or search citation data to get citations. In addition, Crossref has a text mining opt-in program for publishers. The result of this is that some publishers provide URLs for full text content of their articles. The majority of these links are pay-walled, while some are open access. Using any of the various tools for working with Crossref data, you can filter your search to get only articles with full text links, and further to get only articles with full text links that are open access.

The main interfaces for Crossref in R are rcrossref and crminer. Similar interfaces are available in Ruby (serrano) and Python (habanero).

Datacite is similar to Crossref, but focuses on datasets instead of articles. The main interface for Datacite in R is rdatacite.

Biodiversity Heritage Library

The Biodiversity Heritage Library (BHL) houses scans of biodiversity books, and provides web interfaces and APIs to query and fetch those data. They also provide text of the scanned pages. The main R interace to BHL is through rbhl.

[21] https://fatcat.wiki/

[22] http://www.sherpa.ac.uk/romeo/apimanual.php?la=en&fIDnum=%7C&mode=simple

[23] https://core.ac.uk/

[24] https://dissemin.readthedocs.io/en/latest/api.html

Public Library of Science

The Public Library of Science (PLOS) is one of the largest open access only publishers. They as of this writing have published 2.1 million articles. One of the strongs advantages of PLOS is that they provide an API to their Solr instance, which is a very flexible way to search their articles. The main R interace to PLOS is through rplos.

Elsevier/Scopus

Elsevier is one of the largest publishers. Most of their articles are not open access. However, they have a numbrer of advantages if you have access to their articles: they are one of the few publishers to provide machine readable XML (many publishers do have XML versions of articles, but do not provide it); they are one of the few (two) publishers part of Crossref's text and data mining program. The packages fulltext and crminer can be used to access Elsevier articles through Crossref's TDM program. There's an interface to Scopus article search within fulltext.

arXiv/bioRxiv

arXiv and bioRxiv are preprint publishers, the former in existence for many years, and the latter new on the scene. You can access articles from these publishers through fulltext. arXiv does provide a web API that we hook into; bioRxiv does not, but we can get you articles nonetheless.

Pubmed/PMC/Europe PMC

Pubmed/PMC is a corpus/website of NIH funded research in the United States; while Europe PMC is an equivalent for the European Union. You can access articles from Pubmed/PMC through fulltext, and for Europe PMC through europepmc.

Microsoft Academic Research

Microsoft Academic Research (MAR) is a search engine for research articles. You can use their GUI web interface to search, and they provide APIs for programmatic access. The R interface for MAR is microdemic; and fulltext hooks into `microdemic` as well for article search and abstract retrieval.

Directory of Open Access Journals

Directory of Open Access Journals (DOAJ) maintains data on open access journals, as well as some portion of the articles in those journals. Thus, you can search for journals as well as articles with DOAJ. The R interface for DOAJ is jaod.

JSTOR

JSTOR's Data for Research program gives institutions with access to JSTOR, access to full text of articles within JSTOR. There is no way however to make the interaction with JSTOR completely programmatic, thus making reproducible research very difficult. Nonetheless, there is an R package (jstor) for using data from JSTOR's Data for Research.

ORCID

ORCID (https://orcid.org/) is an organization keeping track of identifiers and metadata for researchers around the world. Individuals can optionally maintain metadata on their scholarly works connected to their account with ORCID. Thus, across all of ORCID, a significant cache of metadata is accruing on scholarly works, their funding amounts, collaborators, etc., useful for bibliometrics research and more. The R interface for ORCID is rorcid.

Citoid/Open Citation Corpus

The Open Citation Corpus (http://opencitations.net/) holds records of which articles cite which other articles, allowing for all important research on the scholarly web of citation. Citation data has been very closely guarded until recently, but the largest publishers are still not contributing to the Open Citation Corpus. The R interface to the Open Citation Corpus is rcitoid.

Fatcat

Fatcat is a project from Ben Newbold of the Internet Archive Labs. It is a "versioned, publicly-editable catalog of research publications: journal articles, conference proceedings, pre-prints, blog posts". Fatcat is currently does not have an R client, but is used inside of the fulltext package.

SHERPA/RoMEO

SHERPA/RoMEO (http://sherpa.mimas.ac.uk/romeo/index.php) aggregates and analyses publisher open access policies and provides summaries of self-archiving permissions and conditions of rights given to authors. The [rromeo][] is an R interface to SHERPA/RoMEO.

CORE

CORE (https://core.ac.uk/) touts itself as the world's largest collection of open access research articles, providing metadata on journals and articles, as well as access to the full text of articles. The rcoreoa R package interfaces with the CORE API.

Dissemin

Dissemin (https://dissem.in/) detects papers behind pay-walls and invites authors to upload them to an open repository. Dissemin provides metadata including links to open versions of articles. The [dissemr][] R package interfaces with the Dissemin API.

fulltext: a toolset for text mining in R

fulltext is a general purpose R package for the data part of text mining: search for articles, get links to articles, get article abstracts, and fetch full text of articles. The `fulltext` package is always adding additional data sources as time allows (See Table 1). Starting from searching for articles, the outputs of search can be fed into a function to get links to those articles, or to get abstracts for those articles, or to fetch their full text. The following is a breakdown of the major distinct parts of `fulltext`.

Search

`ft_search()` provides search access to nine different data sources (PLOS, BMC, Crossref, Entrez, arXiv, bioRxiv, Europe PMC, Scopus, Microsoft Academic), creating a mostly unified interface to all data sources. The parts of each data source that are common are for the most part factored out into the parameters of the `ft_search()` function: query term(s), pagination (number of results, result number to start at). In addition, we allow the user to pass on data source specific options to refine the search per data source.

With `ft_search()`, you can query any combination of the nine data sources at once. The returned object is a list, with access to results of each data source by its name (e.g., `$plos`, or `$crossref`). For each data source, the returned object does vary because the returned data from each data source widely varies; for the most part data.frame's are returned. For those data sources not queried, their slot is empty.

One important aspect of the research result we highlight is the licenses in the returned data for each data source.

```
x <- ft_search(query = 'ecology', from = c("plos", "crossref"))
```

The results for this PLOS search have all CC-BY licenses

```
x$plos
#> Query: [ecology]
#> Records found, returned: [47257, 10]
#> License: [CC-BY]
#>                               id
#> 1  10.1371/journal.pone.0001248
#> 2  10.1371/journal.pone.0059813
#> 3  10.1371/journal.pone.0080763
#> 4  10.1371/journal.pone.0155019
#> 5  10.1371/journal.pone.0175014
#> 6  10.1371/journal.pone.0150648
#> 7  10.1371/journal.pone.0208370
#> 8  10.1371/journal.pcbi.1003594
#> 9  10.1371/journal.pone.0102437
#> 10 10.1371/journal.pone.0166559
```

Whereas the results for this Crossref search have mixed licenses

```
x$crossref
#> Query: [ecology]
#> Records found, returned: [164657, 10]
```

```
#> License: [variable, see individual records]
#>    archive                container.title    created  deposited
#> 1  Portico                        Ecology 2006-05-03 2018-08-04
#> 2  Portico                        Ecology 2006-05-03 2018-08-04
#> 3       NA                        Ecology 2006-05-03 2018-08-04
#> 4       NA                        Ecology 2006-05-03 2018-08-04
#> 5       NA                        Ecology 2006-05-03 2018-08-04
#> 6       NA                        Ecology 2006-05-03 2018-08-04
#> 7       NA                        Ecology 2006-05-09 2018-08-01
#> 8  Portico                        Ecology 2017-04-26 2019-03-08
#> 9       NA Trends in Ecology & Evolution 2002-07-25 2017-06-14
#> 10      NA Journal of Industrial Ecology 2014-11-21 2017-06-23
#> Variables not shown: published.print (chr), published.online (chr), doi
#>      (chr), indexed (chr), issn (chr), issue (chr), issued (chr), member
#>      (chr), page (chr), prefix (chr), publisher (chr), reference.count
#>      (chr), score (chr), source (chr), title (chr), type (chr), url (chr),
#>      volume (chr), author (list), link (list), license (list), subject
#>      (chr), alternative.id (chr), subtitle (chr), reference (list)
```

You can dig into the license field for each article, with URLs holding information on each license

```
vapply(x$crossref$data$license, function(w) w$URL[1], "")
#>  [1] "http://doi.wiley.com/10.1002/tdm_license_1.1"
#>  [2] "http://doi.wiley.com/10.1002/tdm_license_1.1"
#>  [3] "http://doi.wiley.com/10.1002/tdm_license_1"
#>  [4] "http://doi.wiley.com/10.1002/tdm_license_1.1"
#>  [5] "http://doi.wiley.com/10.1002/tdm_license_1"
#>  [6] "http://doi.wiley.com/10.1002/tdm_license_1"
#>  [7] "http://doi.wiley.com/10.1002/tdm_license_1"
#>  [8] "http://doi.wiley.com/10.1002/tdm_license_1.1"
#>  [9] "http://www.elsevier.com/tdm/userlicense/1.0/"
#> [10] "http://doi.wiley.com/10.1002/tdm_license_1.1"
```

Links

`ft_links()` provides two pathways to get links (URLs) for articles, with a choice of four different data sources (PLOS, BMC, Crossref, Entrez). First, you can use `ft_search()`, then pass the output of that function to `ft_links()`.

```
out <- ft_search(query = "ecology", from = "entrez")
ft_links(out)
#> <fulltext links>
#> [Found] 6
#> [IDs] ID_30964001 ID_30962485 ID_30962432 ID_30952928 ID_30674747
#>       ID_30674743 ...
```

Second, you can pass DOIs directly to `ft_links()`. Both end up at the same point, links for each article, if they could be found for the user selected data source.

```
# FIXME
ft_links(out$entrez$data$doi)
```

The biggest caveat with `ft_links()` is that we can't gaurantee that the links will work. Link rot is one way in which the links may not work: link rot is when the URL does not point to the original content anymore, or fails altogether. Additionally, with Crossref, publishers can deposit URLs for articles, but they make change the URLs at some later date but not update the URLs with Crossref.

Abstracts

`ft_abstract()` provides access to article abstracts from four different data sources (PLOS, Scopus, Microsoft Academic Research, Crossref). The only way to use the function is to pass article identifiers, which are for the most DOIs.

The advantage of abstracts over full text is that abstracts can often be retrieved even for paywalled articles. That is, you can have much broader coverage of the articles you're targeting relative to full text.

If you are after abstracts, and you are already getting or already have full text, and if the articles are in XML format, then you can use pubchunks to extract out the abstracts.

Fetch full text

`ft_get()` fetchs full text of articles from many different data sources. From the DOIs that are passed in to the function, we detect the publisher, and there are specific plugins for certain publishers: AAAS, American Institute of Physics, American Society of Clinical Oncology, American Society for Microbiology, arXiv, bioRxiv, BiomedCentral, Copernicus, Crossref, Elife, Elsevier, Pubmed/PMC via NCBI's Entrez, Frontiers, IEEE, Informa, Instituto de Investigaciones Filologicas, American Medical Association, Microbiology Society, PeerJ, Pensoft, PLOS, PNAS, Royal Society of Chemistry, ScienceDirect, Scientific Societies, and Wiley.

If there's no built-in plugin for the publisher already, we use the FTDOI API (https://ftdoi.org) to try to get the link for the full text of the article. If the FTDOI API doesn't bear fruit, we search Crossref for a link to the full text. If Crossref doesn't have any full text links, we give up.

Since users can go through a lot of article requests, we cache successfully downloaded articles, and keep that knowledge consistent across R sessions; all subsequent requests for the same article just use the cached version. Additionally, all errors in `ft_get()` are collected in a tidy data.frame in the output of the function to help the user quickly determine what went wrong.

How to text mine from R: Three case studies

Case study 1: Citation mining

In this example, xxxx

Load libraries

```
library("rcrossref")
library("rplos")
library("rorcid")
library("rcitoid")
library("citecorp")
```

rcrossref

Using `rcrossref` for Crossref data:

```
x <- cr_works(query="NSF")
head(x$data)
#> # A tibble: 6 x 32
#>   alternative.id container.title created deposited published.print doi
#>   <chr>          <chr>           <chr>   <chr>     <chr>           <chr>
#> 1 S106352031630~ Applied and Co~ 2016-0~ 2019-02-~ 2018-03        10.1~
#> 2 <NA>           Biogeosciences~ 2017-0~ 2017-07-~ <NA>          10.5~
#> 3 <NA>           Global Biogeoc~ 2018-0~ 2019-01-~ 2018-10        10.1~
#> 4 <NA>           IEEE Communica~ 2016-1~ 2017-12-~ 2017           10.1~
#> 5 S0021782414O0~ Journal de Mat~ 2014-0~ 2018-10-~ 2014-10        10.1~
#> 6 123            Light: Science~ 2019-0~ 2019-01-~ 2019-12        10.1~
#> # ... with 26 more variables: indexed <chr>, issn <chr>, issue <chr>,
#> #   issued <chr>, member <chr>, page <chr>, prefix <chr>, publisher <chr>,
#> #   reference.count <chr>, score <chr>, ...
```

Case study 2: Abstract mining

Sometimes you just need abstracts for your research question. The benefit of only needing abstracts, and not need full text, is that there's many more articles that will have abstracts available than have their full text available.

As an example, let's say you xxxx

```
library("fulltext")
```

xxxxx

Using `fulltext`:

```
res <- ft_search("ecology", from = "crossref",
  crossrefopts = list(filter = c(has_abstract = TRUE)))
ids <- res$crossref$data$doi
out <- ft_abstract(x = ids, from = "crossref")
abstracts <- vapply(out$crossref, "[[", "", "abstract")
```

252 Using quanteda, read the abstracts into a corpus

```
library("quanteda")
corp <- corpus(abstracts)
docvars(corp) <- ids
```

253 Get a summary of the abstracts

```
summary(corp)
#> Corpus consisting of 10 documents:
#>
#>    Text Types Tokens Sentences                    V1
#>   text1   143    262        10   10.2458/v22i1.21112
#>   text2   117    244         6   10.2458/v17i1.21696
#>   text3    75    118         4   10.2458/v25i1.23119
#>   text4     5      8         1    10.2458/v1i1.21154
#>   text5   105    171         7   10.1155/2011/868426
#>   text6   112    181         6   10.1155/2012/273413
#>   text7   117    240         8 10.5194/we-13-91-2013
#>   text8   140    245         9 10.5194/we-13-95-2013
#>   text9   107    202         7   10.1155/2014/198707
#>  text10   118    224         6   10.5402/2011/897578
#>
#> Source: /Users/sckott/github/ropensci/textmine/use-cases/* on x86_64 by sckott
#> Created: Thu Apr 11 11:20:19 2019
#> Notes:
```

254 Use the `kwic()` function to see a word in context across the abstracts

```
kwic(corp, pattern = "ecology")
#>
#>   [text1, 33] knowledge production within critical political | ecology |
#>   [text1, 50]              in scientific articles on dryland | ecology |
```

```
#>  [text1, 204]                 to equilibrium models in range | ecology |
#>  [text1, 246]    communal areas.Keywords: Critical political | ecology |
#>  [text1, 255]                , scientific models, rangeland | ecology |
#>    [text2, 5]                           < jats:p> Political | ecology |
#>   [text2, 23]        manifestations of political economy and | ecology |
#>   [text2, 45]                     I try to extend political | ecology |
#>  [text2, 149]                   , in dialogue with political | ecology |
#>  [text2, 177]            people and resources that political | ecology |
#>  [text2, 229]      indigeneity scholars.Key words: political | ecology |
#>   [text3, 71]                   an analysis from a political | ecology |
#>  [text3, 114]               system, supermarkets, political | ecology |
#>  [text6, 134]               was observed when allopatry and | ecology |
#>  [text7, 167]             ecosystem should be considered for | ecology |
#>  [text7, 185]                       the" four-color issue of | ecology |
#>  [text7, 201]             step toward advancing knowledge in | ecology |
#>  [text9, 195]         or for theoretical studies integrating | ecology |
#>
#>  . This article is a
#>  , and investigates the functions
#>  , and the fence-line photographs
#>  , fence-line photography, scientific
#>  , Southern Africa</
#>  has expanded in multiple new
#>  in the" problem"
#>  to engage with ethnic studies
#>  approaches to better understand the
#>  focuses on cannot be adequately
#>  , coloniality, Maidu,
#>  standpoint allows a different interpretation
#>  </ jats:p>
#>  act together, leading to
```

```
#>  "? Here, I
#>  ", and propose that
#>  and conservation biology. In
#>  and biogeography.</
```

Case study 3: Full text mining

In this example, xxxx

```
library("fulltext")
# library("crminer")
```

Search for articles

Search for the term *ecology* in PLOS journals.

```
(res1 <- ft_search(query = 'ecology', from = 'plos'))
#> Query:
#>   [ecology]
#> Found:
#>   [PLoS: 47272; BMC: 0; Crossref: 0; Entrez: 0; arxiv: 0; biorxiv: 0; Europe PMC: 0; Scopus:
#> Returned:
#>   [PLoS: 10; BMC: 0; Crossref: 0; Entrez: 0; arxiv: 0; biorxiv: 0; Europe PMC: 0; Scopus: 0;
```

Each publisher/search-engine has a slot with metadata and data

```
res1$plos
#> Query: [ecology]
#> Records found, returned: [47272, 10]
#> License: [CC-BY]
#>                              id
#> 1  10.1371/journal.pone.0001248
#> 2  10.1371/journal.pone.0059813
#> 3  10.1371/journal.pone.0080763
```

```
#> 4  10.1371/journal.pone.0155019
#> 5  10.1371/journal.pone.0175014
#> 6  10.1371/journal.pone.0150648
#> 7  10.1371/journal.pone.0208370
#> 8  10.1371/journal.pcbi.1003594
#> 9  10.1371/journal.pone.0102437
#> 10 10.1371/journal.pone.0166559
```

Get full text

Using the results from `ft_search()` we can grab full text of some articles

```
(out <- ft_get(res1))
#> <fulltext text>
#> [Docs] 10
#> [Source] ext - /Users/sckott/Library/Caches/R/fulltext
#> [IDs] 10.1371/journal.pone.0001248 10.1371/journal.pone.0059813
#>       10.1371/journal.pone.0080763 10.1371/journal.pone.0155019
#>       10.1371/journal.pone.0175014 10.1371/journal.pone.0150648
#>       10.1371/journal.pone.0208370 10.1371/journal.pcbi.1003594
#>       10.1371/journal.pone.0102437 10.1371/journal.pone.0166559 ...
```

Extract text from pdfs

Ideally for text mining you have access to XML or other text based formats. However, sometimes you only have access to PDFs. In this case you want to extract text from PDFs. `fulltext` can help with that.

You can extract from any pdf from a file path, like:

```
path <- system.file("examples", "example1.pdf", package = "fulltext")
ft_extract(path)
#> <document>/Library/Frameworks/R.framework/Versions/3.5/Resources/library/fulltext/examples/ex
#>   Title: Suffering and mental health among older people living in nursing homes---a mixed-met
```

```
#>   Producer: pdfTeX-1.40.10
#>   Creation date: 2015-07-17
```

Extract text chunks

Requires the pubchunks library. Here, we'll search for some PLOS articles, then get their full text, then extract various parts of each article with `pub_chunks()`.

```
library("pubchunks")
res <- ft_search(query = "ecology", from = "plos", limit = 3)
x <- ft_get(res)
x %>% ft_collect() %>% pub_chunks(c("doi", "history")) %>% pub_tabularize()
#> $plos
#> $plos$`10.1371/journal.pone.0001248`
#>                            doi history.received history.accepted
#> 1 10.1371/journal.pone.0001248       2007-07-02       2007-11-06
#>   .publisher
#> 1       plos
#>
#> $plos$`10.1371/journal.pone.0059813`
#>                            doi history.received history.accepted
#> 1 10.1371/journal.pone.0059813       2012-09-16       2013-02-19
#>   .publisher
#> 1       plos
#>
#> $plos$`10.1371/journal.pone.0080763`
#>                            doi history.received history.accepted
#> 1 10.1371/journal.pone.0080763       2013-08-15       2013-10-16
#>   .publisher
#> 1       plos
```

Future directions

Text mining will always be a complex task given all the layers involved: often temporal time-span of research questions; varied permissions among researchers and their articles they're trying to access; varied approaches to getting full text (xml vs pdf vs plain text); and more.

Programmatic text mining is a first step towards making text mining easier. The R ecosystem is an especially good place to do text mining because there are many packages for text mining analysis, and endless packages for any required statistical analyses. In addition, rOpenSci and others are building up a set of packages in R for searching for and acquiring full text programatically to help make the research workflow as reproducible as possible.

Future work for `fulltext` includes:

1. Adding more publisher plugins
2. Fine tuned user control over publishers
3. Improve VPN/proxy controls
4. Incorporate more search engines to help resolve URLs for fulltext versions
5. Improve documentation

With respect to what publishers can do to make text mining easier, publishers should:

1. provide XML if they have it
2. not change URL patterns so often, or at all
3. maintain consistent URL patterns among journals, years, etc.
4. keep their Crossref metadata up to date
5. open up their citation data

Acknowledgments

xxxx

Data Accessibility

All scripts and data used in this paper can be found in the permanent data archive Zenodo under the digital object identifier (DOI). This DOI corresponds to a snapshot of the GitHub repository at https://github.com/ropensci/textmine. Software can be found at https://github.com/ropensci/xxx, xxxx, all under MIT licenses.

References

Ba M., Bossy R. 2016. Interoperability of corpus processing work-flow engines: The case of alvisnlp/ml in openminted. In: *Proceedings of the workshop on cross-platform text mining and natural language processing interoperability (interop 2016) at lrec*. 15–18.

Camerer CF., Dreber A., Forsell E., Ho T-H., Huber J., Johannesson M., Kirchler M., Almenberg J., Altmejd A., Chan T., Heikensten E., Holzmeister F., Imai T., Isaksson S., Nave G., Pfeiffer T., Razen M., Wu H. 2016. Evaluating replicability of laboratory experiments in economics. *Science* 351:1433–1436.

Camerer CF., Dreber A., Holzmeister F., Ho T-H., Huber J., Johannesson M., Kirchler M., Nave G., Nosek BA., Pfeiffer T., Altmejd A., Buttrick N., Chan T., Chen Y., Forsell E., Gampa A., Heikensten E., Hummer L., Imai T., Isaksson S., Manfredi D., Rose J., Wagenmakers E-J., Wu H. 2018. Evaluating the replicability of social science experiments in nature and science between 2010 and 2015. *Nature Human Behaviour* 2:637–644.

Cañada A., Capella-Gutierrez S., Rabal O., Oyarzabal J., Valencia A., Krallinger M. 2017. LimTox: A web tool for applied text mining of adverse event and toxicity associations of compounds, drugs and genes. *Nucleic Acids Research* 45:W484–W489.

Chaix E., Deléger L., Bossy R., Nédellec C. 2019. Text mining tools for extracting information about microbial biodiversity in food. *Food Microbiology* 81:63–75.

Ding Z., Li Z., Fan C. 2018. Building energy savings: Analysis of research trends based on text mining. *Automation in Construction* 96:398–410.

Kong X., Gerstein MB. 2018. Text mining systems biology: Turning the microscope back on the observer. *Current Opinion in Systems Biology* 11:117–122.

McCallen E., Knott J., Nunez-Mir G., Taylor B., Jo I., Fei S. 2019. Trends in ecology: Shifts in ecological research themes over the past four decades. *Frontiers in Ecology and the Environment* 17:109–116.

Muñoz G., Kissling WD., Loon EE van. 2019. Biodiversity observations miner: A web application to unlock primary biodiversity data from published literature. *Biodiversity Data Journal* 7.

Open Science Collaboration. 2015. Estimating the reproducibility of psychological science. *Science* 349:aac4716–aac4716.

Piwowar H., Priem J., Larivière V., Alperin JP., Matthias L., Norlander B., Farley A., West J., Haustein S. 2018. The state of OA: A large-scale analysis of the prevalence and impact of open access articles. *PeerJ* 6:e4375.

Sinclair L., Ijaz UZ., Jensen LJ., Coolen MJ., Gubry-Rangin C., Chroňáková A., Oulas A., Pavloudi C., Schnetzer J., Weimann A., Ijaz A., Eiler A., Quince C., Pafilis E. 2016. `Seqenv`: Linking sequences to environments through text mining. *PeerJ* 4:e2690.

Usai A., Pironti M., Mital M., Mejri CA. 2018. Knowledge discovery out of text data: A systematic review via text mining. *Journal of Knowledge Management* 22:1471–1488.